

# Class 8: Breast Cancer Analysis Mini Project

Serena Quezada (PID: A18556865)

## Table of contents

Background . . . . .	1
Data Import . . . . .	1
1. Data Exploration . . . . .	3
2. Principal Component Analysis (PCA) . . . . .	4
PCA screeplot . . . . .	8
Communicating PCA results . . . . .	9
3. Hierarchical clustering . . . . .	10
5. Combining methods (PCA and CLustering) . . . . .	11
7. Prediction . . . . .	14

## Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: “Robust Linear Programming Discrimination of Two Linearly Inseparable Sets”.

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

## Data Import

Data was downloaded from the class website as a CSV file

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names = 1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001		0.14710
842517	0.08474	0.07864	0.0869		0.07017
84300903	0.10960	0.15990	0.1974		0.12790
84348301	0.14250	0.28390	0.2414		0.10520
84358402	0.10030	0.13280	0.1980		0.10430
843786	0.12780	0.17000	0.1578		0.08089
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
84358402	0.01756		0.005115	22.54	16.67
843786	0.02165		0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622		0.6656
842517	158.80	1956.0	0.1238		0.1866
84300903	152.50	1709.0	0.1444		0.4245
84348301	98.87	567.7	0.2098		0.8663
84358402	152.20	1575.0	0.1374		0.2050
843786	103.40	741.6	0.1791		0.5249
	concavity_worst	concave.points_worst	symmetry_worst		

842302	0.7119	0.2654	0.4601
842517	0.2416	0.1860	0.2750
84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985
	fractal_dimension_worst		
842302		0.11890	
842517		0.08902	
84300903		0.08758	
84348301		0.17300	
84358402		0.07678	
843786		0.12440	

## 1. Data Exploration

The first column **diagnosis** is the expert opinion on the sample (i.e. patient FNA)

```
wisc.data <- wisc.df[,-1]
dim(wisc.data)
```

[1] 569 30

Store the diagnosis as a vector for use later when we compare our results to those from experts in the field

```
diagnosis <- factor(wisc.df$diagnosis)
```

Q1. How many observations are in this dataset?

There are 569 observations/patients in the dataset

Q2. How many of the observations have a malignant diagnosis?

There are 212 observations that have a malignant diagnosis

```
table(wisc.df$diagnosis)
```

B	M
357	212

Q3. How many variables/features in the data are suffixed with `_mean`?

There are 10 variables/features in the data set that the “`_mean`” suffix

```
#colnames(wisc.data)
length( grep("_mean", colnames(wisc.data)) )
```

```
[1] 10
```

## 2. Principal Component Analysis (PCA)

The `prcomp()` function to do PCA has a `scale=FALSE` default. In general we nearly always want to set this to TRUE so our analysis is not dominated by columns/variables in our dataset that have high standard deviation and mean when compared to others just because the units of measurement are on different units/scales.

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

It is 0.4427


Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

The cumulative exceeds 0.70 after the third component (0.72636).

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

The cumulative proportion first passes the 0.90 threshold after the seventh component .91010.

```
biplot(wisc.pr,
       xlab = as.character(diagnosis),
       cex = 0.6)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

It is difficult to understand because it has lots of plots and numbers bunched up together and you aren't able to read anything or see the plots.


Generating a more standard scatter plot

```
plot(wisc.pr$x[,1:2],  
      col = diagnosis,  
      xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components and 3. What do you notice about these plots?

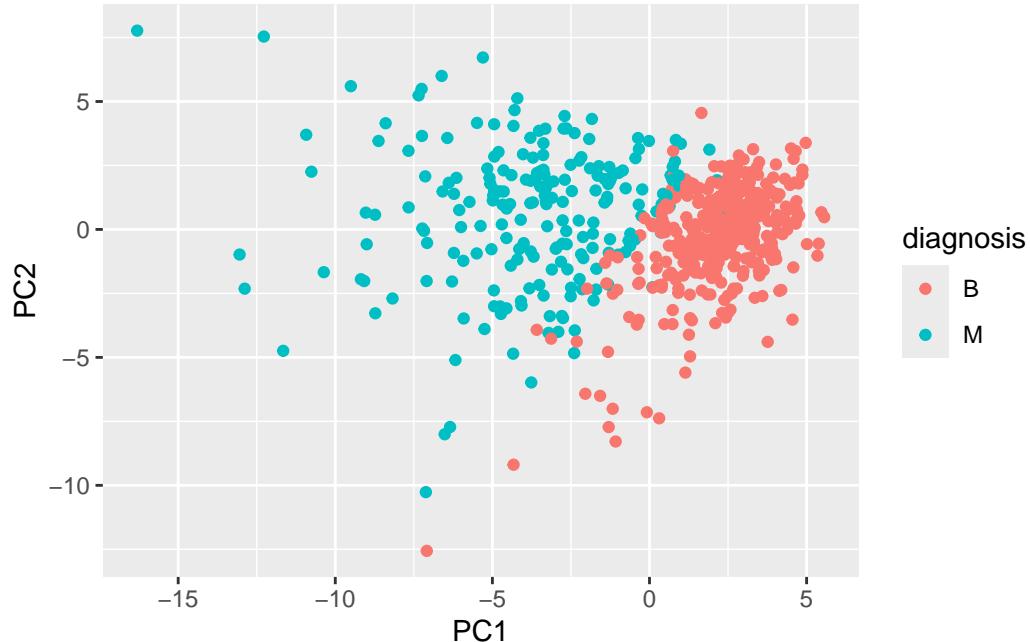
```
plot(wisc.pr$x[, c(1,3)],  
      col = diagnosis,  
      xlab = "PC1", ylab = "PC3")
```



The first plot is very different from the 2nd and 3rd plot where in the 1st plot everything was bunched up together and plots couldn't be seen. In plot 2 and 3 it explains more variance, and is capturing a separation between the red and black samples properly.

```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



### PCA screeplot


A plot of how much variance each PC captures. We can get this from `wisc.pr$sdev` or from the output of `summary(wisc.pr)`

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pc.var <- wisc.pr$sdev^2

pve <- pc.var / sum(pc.var)
plot(pve, xlab = "Principle Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0,1), type = "o")
```



## Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation["concave.points_mean", 1]
```

[1] -0.2608538

```
wisc.pr$rotation["concave.points_mean", "PC1"]
```

[1] -0.2608538

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

The minimum number is 5 principal components to explain 80% of the variance data, because PC5 exceeds 80% threshold.


### 3. Hierarchical clustering

```
data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled, method = "euclidean")

wisc.hclust <- hclust(data.dist, method = "complete")

plot(wisc.hclust)
abline(h=150, col = "red", lty = 2)
```



```
data.dist  
hclust (*, "complete")
```

```
h4 <- wisc.hclust$height[length(wisc.hclust$height) - 4 +1]  
h4
```

```
[1] 18.63658
```

```
clusters_at_h <- cutree(wisc.hclust, h=18.63658)  
length(unique(clusters_at_h))
```

```
[1] 5
```

## 5. Combining methods (PCA and CLustering)

Clustering the original data was not very productive. The PCA results looked promising. Here we combine these methods by clustering from our PCA results. In other words "clustering in PC space"...


```
grps <- cutree(wisc.hclust, k = 2)
table(grps)
```

```
grps
 1  2
567  2
```


```
table(grps, diagnosis)
```

```
diagnosis
grps   B   M
 1 357 210
 2   0   2
```

```
plot(wisc.pr$x[, 1:2], col = grps)
```



```
plot(wisc.pr$x[, 1:2], col = diagnosis)
```



```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
dist7 <- dist(wisc.pr$x[, 1:7], method = "euclidean")

wisc.pr.hclust <- hclust(dist7, method = "ward.D2")

wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k = 2)
table(wisc.pr.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.pr.hclust.clusters   B   M
      1 28 188
      2 329 24

```

The model does a good job of separating the two diagnoses. The model captures most cancers while keeping false-positive alerts for healthy patients relatively low.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```

wisc.df    <- read.csv("WisconsinCancer.csv", row.names = 1)
diagnosis <- factor(wisc.df$diagnosis)                      # ensure it is a factor

wisc.data <- wisc.df[ , setdiff(names(wisc.df), "diagnosis")]

wisc.km <- kmeans(wisc.data, centers = 4)

dist.mat    <- dist(scale(wisc.data), method = "euclidean")
wisc.hclust <- hclust(dist.mat, method = "complete")
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)

table(wisc.km$cluster, diagnosis)

```

```

diagnosis
      B   M
1 94 87
2 1 100
3 0 19
4 262 6

```

```
table(wisc.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.hclust.clusters   B   M
      1 12 165
      2 2 5
      3 343 40
      4 0 2

```

The k-means model places each patient into one of the 4 clusters and each cluster is made up of a single diagnosis, ex. cluster 1 contains only malignant, cluster 2 contains benign and so on. This model is excellent at avoiding false-positive alerts which is good when you want to avoid unnecessary follow-up for benign patients. The hierarchical clustering model merges the 2 groups that cause the smallest increase in within-cluster variance.

## 7. Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			


```
col.vec <- ifelse(diagnosis == "M", "red", "black")

plot(wisc.pr$x[, 1:2],
      col = col.vec,
      xlab = "PC 1", ylab = "PC 2",
      main = "PCA scores coloured by diagnosis")

points(npc[, 1], npc[, 2], col = "blue", pch = 16, cex = 3)

text(npc[, 1], npc[, 2], labels = c(1, 2), col = "white")
```

### PCA scores coloured by diagnosis



Q18. Which of these new patients should we prioritize for follow up based on your results?

The patients that fall into the clusters that are dominated by malignant diagnoses should be the first ones to schedule for a follow up.