

Class 9 structural bioinformatics part 1

Serena Quezada (PID: A18556865)

The main database for structural biology is called the PDB. Let's have a look at what it contains:

1. Introduction to the RCSB Protein

```
read.csv("Data Export Summary.csv")
```

	Molecular.Type	X.ray	EM	NMR	Integrative	Multiple.methods
1	Protein (only)	176,204	20,299	12,708	342	218
2	Protein/Oligosaccharide	10,279	3,385	34	8	11
3	Protein/NA	9,007	5,897	287	24	7
4	Nucleic acid (only)	3,066	200	1,553	2	15
5	Other	173	13	33	3	0
6	Oligosaccharide (only)	11	0	6	0	1
	Neutron	Other	Total			
1	83	32	209,886			
2	1	0	13,718			
3	0	0	15,222			
4	3	1	4,840			
5	0	0	222			
6	0	4	22			

```
library(readr)
stats <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 9
-- Column specification -----
Delimiter: ","
chr (1): Molecular Type
```

```
dbl (4): Integrative, Multiple methods, Neutron, Other
num (4): X-ray, EM, NMR, Total

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
stats
```

```
# A tibble: 6 x 9
  `Molecular Type` `X-ray`    EM    NMR Integrative `Multiple methods` Neutron
  <chr>           <dbl>    <dbl>  <dbl>        <dbl>          <dbl>    <dbl>
1 Protein (only)   176204  20299 12708        342        218      83
2 Protein/Oligosacch~ 10279   3385   34         8        11       1
3 Protein/NA        9007    5897   287        24        7       0
4 Nucleic acid (only) 3066    200   1553        2        15      3
5 Other              173     13    33         3        0       0
6 Oligosaccharide (o~ 11      0     6          0        1       0
# i 2 more variables: Other <dbl>, Total <dbl>
```

```
n.total <- sum(stats$Total, na.rm = TRUE)
```

```
stats$Total
```

```
[1] 209886 13718 15222 4840    222     22
```

Q1. What percentage of structures in the PDB are solved by X-ray and Electron Microscopy

```
n.xray <- sum(stats$`X-ray`)
n.xray / n.total * 100
```

```
[1] 81.48087
```

Q2. What proportion of structures in the PDB are protein?

```
round(stats$Total[1]/n.total * 100,2)
```

```
[1] 86.05
```

Q3. Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB

There are about 4,865 HIV-protease structures

2. Visualizing the HIV-1 protease structure

Package for structural bioinformatics

```
library(bio3d)  
  
hiv <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
hiv
```

Call: read.pdb(file = "1hsg")

Total Models#: 1
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [HOH (127), MK1 (1)]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
       calpha, remark, call
```

Let's first use the Mol* viewer to explore this structure



> Q4. Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?

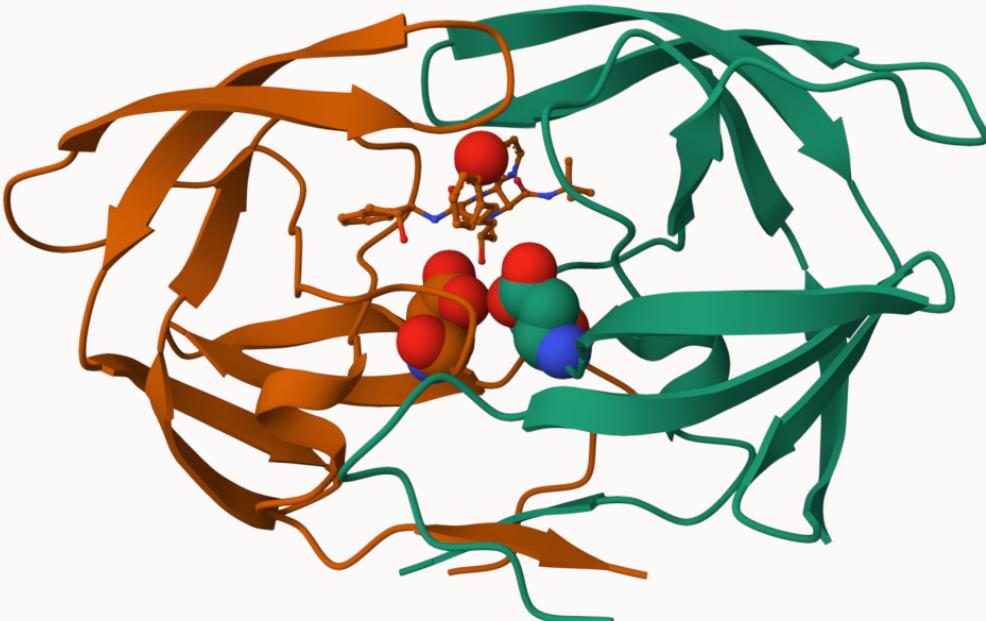
We only see one atom per water molecule in this structure because its a simplified symbol representing the entire molecule, not a single real atom.

Q5. There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have?

The residue number this water molecule has is HOH 306

3. Introduction to Bio3D in R

Q6. And a view of the ligand with catlytic ASP 25 amino-acids (spacefill) and the all important active site water molecule (spacefill): Can you think of a way in which indinavir, or even larger ligands and substrates, could enter the binding site?



```
## Reading PDB file data in R
```

```
library(bio3d)  
  
pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
Warning in get.pdb(file, path = tempdir(), verbose = FALSE):  
/var/folders/zk/6hldzf5n74scx33n9zx15450000gn/T//Rtmp64DTAV/1hsg.pdb exists.  
Skipping download
```

```
pdb
```

```
Call: read.pdb(file = "1hsg")
```

```
Total Models#: 1  
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)
```

```
Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,
       calpha, remark, call
```

Q7. How many amino acid residues are there in this pdb object?

There are 198 amino acid residues

Q8. Name one of the two non-protein residues?

HOH, which is a water molecule

Q9. How many protein chains are in this structure?

There are 2 protein chains, labeled chain A and B

Quick PDB visualization in R

```
attributes(pdb)
```

```
$names
[1] "atom"     "xyz"      "seqres"    "helix"    "sheet"    "calpha"   "remark"   "call"
$class
[1] "pdb"      "sse"
```

```
head(pdb$atom)
```

```

type eleno elety alt resid chain resno insert      x      y      z o      b
1 ATOM    1      N <NA>   PRO     A     1  <NA> 29.361 39.686 5.862 1 38.10
2 ATOM    2      CA <NA>   PRO     A     1  <NA> 30.307 38.663 5.319 1 40.62
3 ATOM    3      C <NA>    PRO     A     1  <NA> 29.760 38.071 4.022 1 42.64
4 ATOM    4      O <NA>    PRO     A     1  <NA> 28.600 38.302 3.676 1 43.40
5 ATOM    5      CB <NA>   PRO     A     1  <NA> 30.508 37.541 6.342 1 37.87
6 ATOM    6      CG <NA>   PRO     A     1  <NA> 29.296 37.591 7.162 1 38.40
segid elesy charge
1 <NA>      N <NA>
2 <NA>      C <NA>
3 <NA>      C <NA>
4 <NA>      O <NA>
5 <NA>      C <NA>
6 <NA>      C <NA>

```

I can interactively view these PDB objects in R with the new **bio3dview** package. This is not yet on CRAN

To install this I can setup **pak** package and use it to install **bio3dview** from GitHub. In my console I first run

```

install.packages("pak") pak::pak("bioboot/bio3dview")
library(bio3dview) library(NGLVieweR)
view.pdb(hiv)

```

Due to the interactive photo it is not visible on pdf but below is the code that would be used to run the interactive photo

In order to change some settings

```

sel <- atom.select(hiv, resno= 25)
view.pdb(hiv, highlight = sel, colorScheme = "chain", col = c("blue", "orange"), background-
Color ="pink")

```

Prediction protein flexibility

We can run bioinformatics calculation to predict protein dynamics - i.e. functional motions.

We will use the **nmac()** function:

```
adk <- read.pdb("6s36")
```

```
Note: Accessing on-line PDB file
      PDB has ALT records, taking A only, rm.alt=TRUE
```

```
adk
```

```
Call: read.pdb(file = "6s36")

Total Models#: 1
  Total Atoms#: 1898,  XYZs#: 5694  Chains#: 1  (values: A)

  Protein Atoms#: 1654  (residues/Calpha atoms#: 214)
  Nucleic acid Atoms#: 0  (residues/phosphate atoms#: 0)

  Non-protein/nucleic Atoms#: 244  (residues: 244)
  Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]

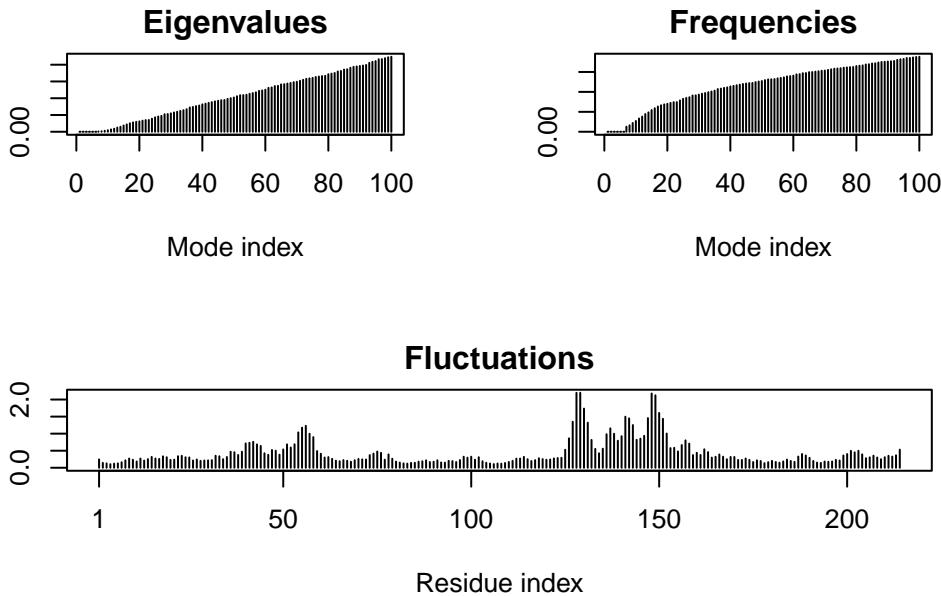
Protein sequence:
  MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGMLRAAVKSGSELGKQAKDIDMAGKLVT
  DELVIALVKERIAQEDCRNGFLLDGFPRТИPQADAMKEAGINVVDYVLEFDVPDELIVDKI
  VGRRVHAPSGRVYHVKNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTAPLIG
  YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG

+ attr: atom, xyz, seqres, helix, sheet,
  calpha, remark, call
```

```
m <- nma(adk)
```

```
Building Hessian...      Done in 0.019 seconds.
Diagonalizing Hessian...  Done in 0.719 seconds.
```

```
plot(m)
```



Generate a “trajectory” of predicted motion

```
mktrj(m, file = "ADK_nma.pdb")
```

```
view.nma(m)
```

4. Comparative structure analysis of Adenylate Kinase

Install packages in the R console NOT your Rmd/Quarto file

```
install.packages("bio3d") install.packages("NGLVieweR")
```

```
install.packages("pak") pak:::pak("bioboot/bio3dview")
```

```
install.packages("BiocManager") BiocManager::install("msa")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

msa if found only on BioConductor and not CRAN

Q11. Which of the above packages is not found on BioConductor or CRAN?

Bio3dview is not found on BioConductor or CRAN

Q12. True or False? Functions from the pak package can be used to install packages from GitHub and BitBucket?

True, the `pak` package is a modern R package manager that can install packages from multiple sources - including CRAN, BioConductor, GitHub and BitBucket